

Pattern Non–Redundancy

We now introduce a judgment expressing that *no branch of a pattern list is redundant*, i.e. each pattern contributes new, previously uncovered values.

We write a sequence of patterns as

$$\vec{p} = (p_1, \dots, p_n).$$

Recall the (partial) matching function

$$\mathbf{match}(p, v) = \theta$$

which is defined precisely when p matches value v and returns a binding environment θ .

Non–Redundancy Judgment

We define a judgment

$$\Delta \vdash \mathbf{NR}(T, \vec{p})$$

read: *under type context Δ , the pattern sequence \vec{p} is non–redundant for type T .*

Intuitively, each pattern p_k must match at least one value of type T that is not matched by any earlier pattern p_1, \dots, p_{k-1} .

We define \mathbf{NR} inductively.

Empty Sequence The empty sequence is trivially non–redundant:

$$[\mathbf{NR}\text{-EMPTY}] \Delta \vdash \mathbf{NR}(T, ())$$

Extending a Non–Redundant Sequence Suppose we already have a non–redundant sequence (p_1, \dots, p_{k-1}) . We may extend it with p_k provided that there exists a value of type T which matches p_k but does *not* match any of the earlier patterns.

Formally:

$$[\mathbf{NR}\text{-CONS}] \Delta \vdash \mathbf{NR}(T, (p_1, \dots, p_{k-1})) \exists v : T. \exists \theta. \mathbf{match}(p_k, v) = \theta \wedge \forall i < k. \mathbf{match}(p_i, v) \text{ is undefined} \Delta \vdash$$

Thus, each new pattern is required to introduce at least one *fresh* value of type T that was not already covered.

Typing Rule with Exhaustiveness and Non–Redundancy

Combining the earlier exhaustiveness judgment \mathbf{Exh} with non–redundancy, we may state a strengthened typing rule for pattern matching.

$$[\mathbf{T}\text{-MATCH-TOTAL}] \Gamma \vdash e : T \forall k. \Gamma, \Gamma_{p_k} \vdash e_k : U \Delta \vdash \mathbf{Exh}(T, \{p_1, \dots, p_m\}) \Delta \vdash \mathbf{NR}(T, (p_1, \dots, p_m)) \Gamma \vdash \mathbf{match}$$

Here:

- Exh guarantees that *every* value of type T is matched by *some* pattern.
- NR guarantees that *each* branch matches *at least one* value that no earlier branch matches.

Non–Redundancy Lemma

We may now formulate a meta–theoretic property of the non–redundancy judgment.

Lemma (Branch Non–Redundancy). Let T be a type and $\vec{p} = (p_1, \dots, p_m)$ a sequence of patterns such that

$$\Delta \vdash \text{NR}(T, \vec{p}).$$

Then for each k with $1 \leq k \leq m$ there exists a value $v_k : T$ and a binding environment θ_k such that:

$$\text{match}(p_k, v_k) = \theta_k \quad \text{and} \quad \forall i < k. \text{match}(p_i, v_k) \text{ is undefined.}$$

In other words, every branch $p_k \Rightarrow e_k$ is *semantically reachable*: there is at least one value of type T for which operational evaluation of the corresponding `match` expression selects that branch.

Sketch of Proof. The proof proceeds by induction on the derivation of $\Delta \vdash \text{NR}(T, \vec{p})$.

- In case NR-EMPTY , the statement is vacuously true.
- In case NR-CONS , the induction hypothesis yields the required witnesses v_i for all $i < k$, and the side condition of the rule provides the witness v_k for the new pattern.

Thus, $\text{NR}(T, \vec{p})$ precisely captures the intended notion of branch non–redundancy in pattern matching.